

# Distributed Linear Regression

Main References

Ameet Talwalkar and Henry Chai, **ML with Large Datasets**, CMU


# Outline

- Machine learning overview
- Large datasets and parallel computing
- Linear regression
- Distributed linear regression

# What is machine learning?

- “Machine learning as a field of study that gives computers **the ability to learn without being explicitly programmed**” - Arthur Samuel (1959)
- “A computer program is said to **learn** from **experience  $E$**  with respect to some class of **tasks  $T$**  and **performance measure  $P$**  if its performance at tasks in  $T$ , as measured by  $P$ , **improves** with experience  $E$ .” - Tom M. Mitchell (1997)
- Machine learning is about **designing algorithms** that **learn from data**.

# Taxonomy of Machine Learning



Labeled Data

Reward

Unlabeled Data

Supervised Learning

Reinforcement Learning



Unsupervised Learning

Quantitative Response

Categorical Response

Regression



Classification



Alpha Go

Dimensionality Reduction

Clustering

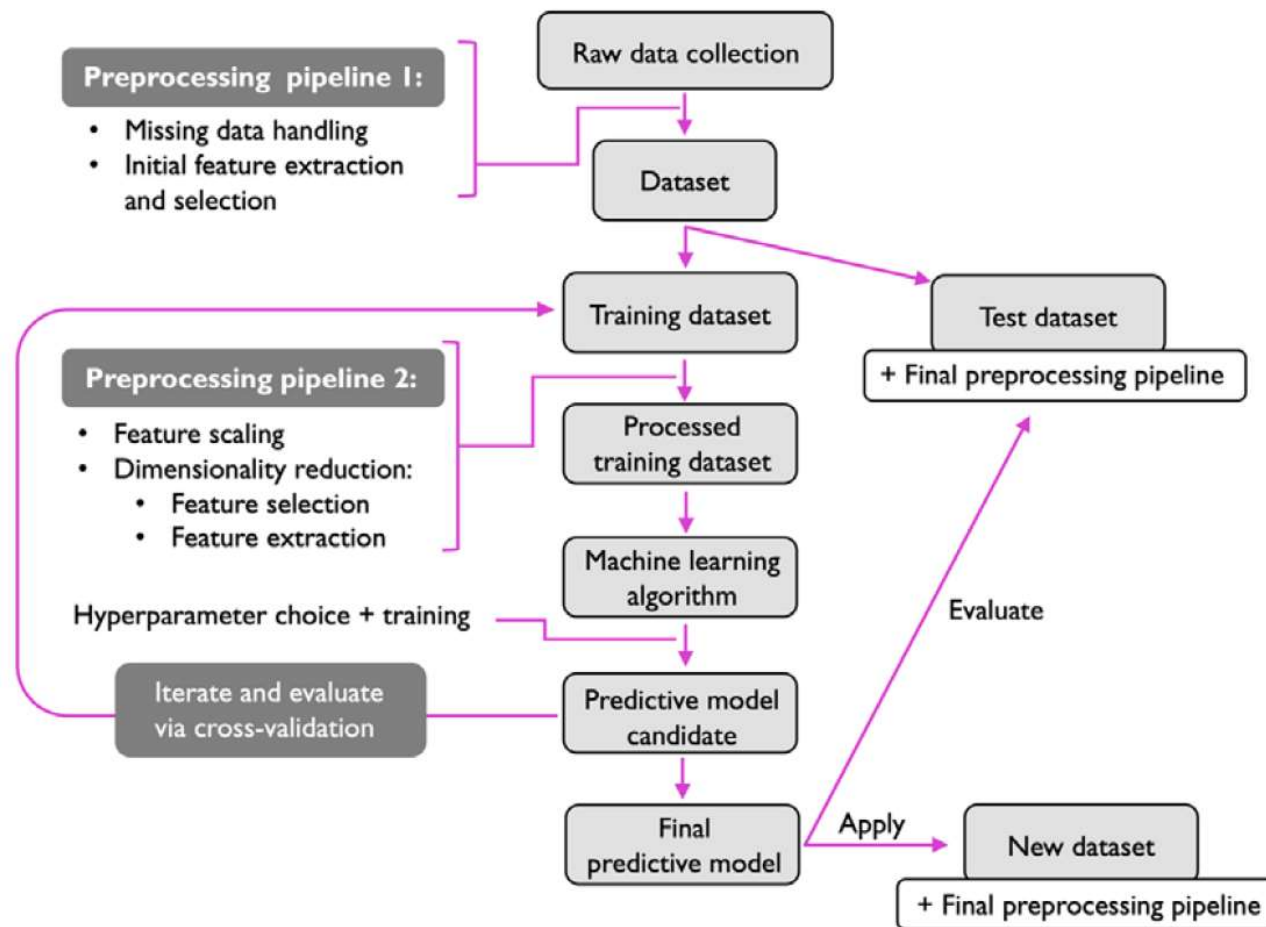


Source: Joseph E. Gonzalez, *AI-Systems Big Ideas*, 2019.

# Machine learning terminology

- **Training example**: A row in a table representing the dataset and synonymous with an observation, record, instance
- **Training**: Model fitting
- **Feature**, abbrev. **x**: A column in a data table or data matrix. Synonymous with predictor, variable, input, attribute, or covariate.
- **Target**, abbrev. **y**: Synonymous with outcome, output, response variable, dependent variable, (class) label, and ground truth.
- **Loss function**: measured for a single data point. Sometimes, also called a **error function**
- **Cost function**: The loss (average or summed) over the entire dataset


# Machine learning workflow/pipeline



Sebastian Raschka, et al., *Giving Computers the Ability to Learn from Data*, In *Machine Learning with PyTorch and Scikit-Learn* (pp. 1–18), Packt Publishing, 2022.

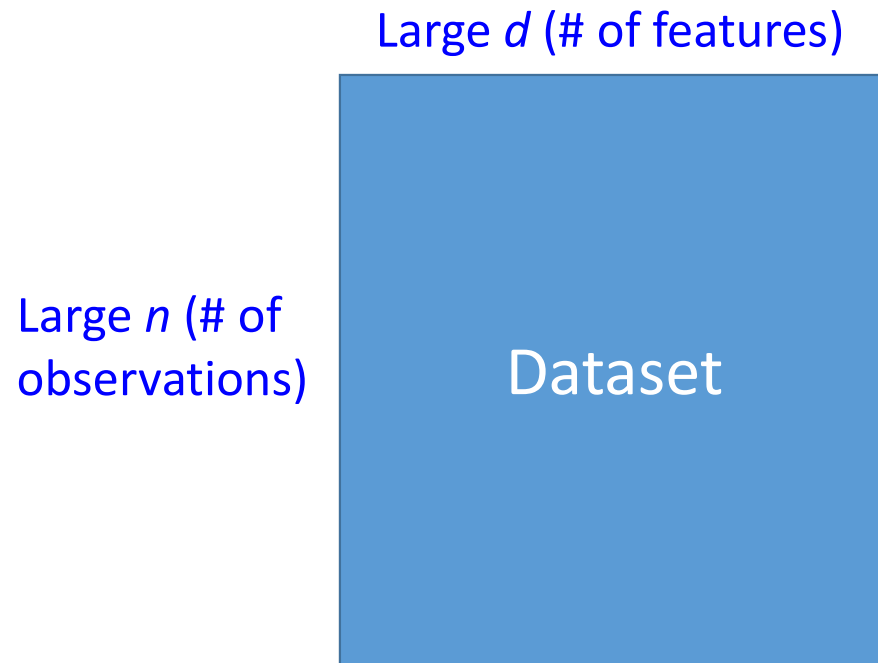
# Model selection and model evaluation

- Suppose we want to compare multiple hyperparameter settings  $\lambda_1, \lambda_2, \dots, \lambda_k$
- For  $i = 1, 2, \dots, k$ 
  - Train a model on  $D_{train}$  using  $\lambda_i$
- Evaluate each model on  $D_{val}$  and find the best hyperparameter setting,  $\lambda_{i^*}$
- Compute the error of a model trained with  $\lambda_{i^*}$  on  $D_{test}$



# Large Datasets

- Datasets can be large in two ways





# Large Datasets: Example

- Image processing
  - Large **n**: potentially massive number of observations (e.g., pictures on the internet)
  - Use-cases: object recognition, annotation generation
- Medical data
  - Large **d**: potentially massive feature set (e.g., genome sequence, electronic medical records, etc...)
  - Use-cases: personalized medicine, diagnosis prediction
- Business analytics
  - Large **n** (e.g., all customers & all products) and **d** (e.g., customer data, product specifications, transaction records, etc...)
  - Use-cases: product recommendations, customer segmentation

# Large $d$ (# of features)

- High-dimensional datasets present numerous issues
  - Curse of dimensionality
  - Overfitting
  - Computational issues
- Strategies
  - Learn low-dimensional representations
  - Perform feature selection

## Large $n$ (# of observations)

- Typically, we consider **exponential** time complexity (e.g.,  $O(2^n)$ ) bad and **polynomial** complexity (e.g.,  $O(n^3)$ ) good
- However, if  **$n$  is massive**, then even  **$O(n)$  can be problematic!**
- Strategies
  - Speed up processing e.g., stochastic gradient descent vs. gradient descent
  - Sampling (make approximations)
  - **Parallel computing (large scale data processing)**

# Parallel computing

- Multi-core processing – scale up
  - Data can fit on one machine
  - Requires high-end (expensive) hardware
  - Simpler algorithms that don't necessarily scale well
- Distributed processing – scale out
  - Data stored across multiple machines
  - Scales to massive problems on commodity (inexpensive) hardware
  - Added complexity of network communication

# Big O Notation

- Used to describe an algorithm's time or space (storage) complexity in terms of the input size

$$f(n) = O(g(n)) \Leftrightarrow \exists c, n_0: f(n) \leq cg(n) \forall n \geq n_0$$

- $O(1)$  = constant time/space
  - $O(\log(n))$  = logarithmic time/space
  - $O(n)$  = linear time/space
- An algorithm's time and space complexity can be different
    - Ex: multiplying an  $m \times n$  matrix with an  $n \times p$  matrix takes  $O(mnp)$  time but the result uses  $O(mn + np + mp)$  storage

# Empirical Risk Minimization – ERM

- ERM is a common framework for supervised learning
- Given:
  - some labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
  - a loss function  $l: Y \times Y \rightarrow R$
  - a hypothesis class or set of functions  $F$
- The goal is to find

$$\hat{f} = \underset{f \in F}{\operatorname{argmin}} \sum_{i=1}^n l(f(x^{(i)}), y^{(i)})$$

with the hope that

$$E_{p(x,y)}[l(f(x), y)] \approx \frac{1}{n} \sum_{i=1}^n l(f(x^{(i)}), y^{(i)})$$

# Empirical Risk Minimization – ERM

- ERM is a common framework for supervised learning
- Given:
  - some labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
  - a loss function  $l: Y \times Y \rightarrow R$
  - a hypothesis class or set of functions  $F$

- The goal is to find

$$\hat{f} = \underset{f \in F}{\operatorname{argmin}} \sum_{i=1}^n l(f(x^{(i)}), y^{(i)})$$

- Depending on the choice of  $F$  and  $l$ , this objective function may be convex (easy to optimize) or non-convex (hard)
- We need to solve this problem for large  $n$  and/or  $d$

# Regression

- Regression is a type of supervised learning
- Given:
  - a labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
  - a loss function  $l: Y \times Y \rightarrow R$ , where  $Y = R$
  - a hypothesis class or set of functions  $F$

- The goal is to find

$$\hat{f} = \underset{f \in F}{\operatorname{argmin}} \sum_{i=1}^n l(f(x^{(i)}), y^{(i)})$$

- Depending on the choice of  $F$  and  $l$ , this objective function may be convex (easy to optimize) or non-convex (hard)
- We need to solve this problem for large  $n$  and/or  $d$



# Linear Regression (Ordinary Least Squares)

- Linear regression is a simplest type of regression
- Given:
  - a labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
  - a loss function  $l(y, y') = (y - y')^2$
  - $F$  = all functions of the form  $f(x) = \theta_0 + \sum_{i=1}^k \theta_i x_i = \theta^T x$
- The goal is to find

$$\hat{f} = \underset{f \in F}{\operatorname{argmin}} \sum_{i=1}^n (\theta^T x^{(i)} - y^{(i)})^2$$

or

$$\hat{\theta} = \underset{\theta}{\operatorname{argmin}} \|X\theta - y\|^2$$

where:

$$X = \begin{bmatrix} x^{(1)T} \\ \vdots \\ x^{(n)T} \end{bmatrix}, y = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix}, x^{(i)} = \begin{bmatrix} 1 \\ x_1^{(i)} \\ \vdots \\ x_d^{(i)} \end{bmatrix}, \theta = \begin{bmatrix} 1 \\ \theta_1 \\ \vdots \\ \theta_d \end{bmatrix}$$

# Linear Regression (Ordinary Least Squares)

$$L_D(\theta) = \|X\theta - y\|^2 = (X\theta - y)^T (X\theta - y)$$

$\vdots$

$$\hat{\theta} = (X^T X)^{-1} X^T y$$

# Regularized Linear Regression

- Regularized empirical risk minimization that penalizes model complexity to deal with overfitting
- Given:
  - some labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
  - a loss function  $l: Y \times Y \rightarrow R$ , where  $Y = R$
  - a hypothesis class or set of functions  $F$
  - a regularizer function  $r: \theta \rightarrow R$
  - a regularization parameter  $\lambda$
- The goal is to find

$$\hat{\theta} = \underset{\theta}{\operatorname{argmin}} \sum_{i=1}^n l(f_{\theta}(x^{(i)}), y^{(i)}) + \lambda r(\theta)$$

# Regularized Linear Regression

$$L_D(\theta) = \|X\theta - y\|^2 + \lambda\|\theta\|^2 = (X\theta - y)^T(X\theta - y) + \lambda\theta^T\theta$$

$\vdots$

$$\hat{\theta} = (X^T X + \lambda I_{d+1})^{-1} X^T y$$

$I_{d+1}$ : is the  $(d + 1) \times (d + 1)$  identity matrix

# Linear Regression – Large $n$ , small $d$

$$\hat{\theta} = (X^T X)^{-1} X^T y$$

- Time complexity for  $(X^T X)^{-1}$  is  $O(d^3)$
- Space complexity for  $X^T X$  is  $O(d^2)$
- Assume  $O(d^3)$  computation and  $O(d^2)$  storage is possible on a single machine
  - We can store and invert  $X^T X$
  - We cannot compute  $X^T X$  on a single machine
  - We cannot store  $X$  on a single machine
- Idea: distribute storage of  $X$  and computation of  $X^T X$ 
  - Store the rows of  $X$  across different machines
  - Compute  $X^T X$  as the sum of outer products

# Matrix Multiplication via Outer Products

$$\begin{bmatrix} a_{11} & \cdots & a_{1m} \\ \vdots & \ddots & \vdots \\ a_{n1} & \cdots & a_{nm} \end{bmatrix} \begin{bmatrix} b_{11} & \cdots & b_{1k} \\ \vdots & \ddots & \vdots \\ b_{m1} & \cdots & b_{mk} \end{bmatrix} = \begin{bmatrix} \sum_{i=1}^m a_{1i}b_{i1} & \cdots & \sum_{i=1}^m a_{1i}b_{ik} \\ \vdots & \ddots & \vdots \\ \sum_{i=1}^m a_{ni}b_{i1} & \cdots & \sum_{i=1}^m a_{ni}b_{ik} \end{bmatrix}$$
$$= \sum_{i=1}^m \begin{bmatrix} a_{1i}b_{i1} & \cdots & a_{1i}b_{ik} \\ \vdots & \ddots & \vdots \\ a_{ni}b_{i1} & \cdots & a_{ni}b_{ik} \end{bmatrix}$$

## Distributed Computation of $(X^T X)^{-1}$

$$X^T X = \begin{bmatrix} \uparrow & \uparrow & \uparrow \\ x^{(1)} & \cdots & x^{(n)} \\ \downarrow & \downarrow & \downarrow \end{bmatrix} \begin{bmatrix} \leftarrow & x^{(1)} & \rightarrow \\ \vdots & \vdots & \vdots \\ \leftarrow & x^{(n)} & \rightarrow \end{bmatrix} = \sum_{i=1}^n x^{(i)} x^{(i)T}$$

**Idea:** distribute  $x^{(i)}$  and compute summands in parallel

# Distributed Computation of $(X^T X)^{-1}$

<b>Worker</b>	$\begin{bmatrix} \leftarrow & x^{(1)T} & \rightarrow \\ \leftarrow & x^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(2)T} & \rightarrow \\ \leftarrow & x^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(5)T} & \rightarrow \\ \leftarrow & x^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nd)$ distributed storage (total)	
<b>Map</b>	$x^{(i)} x^{(i)T}$	$x^{(i)} x^{(i)T}$	$x^{(i)} x^{(i)T}$	$O(nd^2)$ distributed work (total)	$O(d^2)$ local storage
<b>Reduce</b>	$\left( \sum_{i=1}^n x^{(i)} x^{(i)T} \right)^{-1}$			$O(d^3)$ local work	$O(d^2)$ local storage



## Linear Regression – Large $n$ , large $d$

$$\hat{\theta} = (X^T X)^{-1} X^T y$$

- Time complexity for  $(X^T X)^{-1}$  is  $O(d^3)$
- Space complexity for  $X^T X$  is  $O(d^2)$
- Assume  $O(d^3)$  computation and  $O(d^2)$  storage is possible on a single machine
  - We **cannot** store and invert  $X^T X$
  - We **cannot** compute  $X^T X$  on a single machine
  - We **cannot** store  $X$  on a single machine
- **Idea**: Use distributed version of gradient descent

# Gradient Descent

- We're trying to minimize some function  $L$
- Suppose at iteration  $t$ : we're get  $\theta^{(t)}$
- With learning rate  $\eta$ , we update  $\theta^{(t)}$  (at iteration  $t + 1$ ) as follow

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla_{\theta} L(\theta^{(t)})$$

# Gradient Descent for Linear Regression

- Dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$

1. Initialize  $\theta^{(0)} = 0$  (zero vector) and set  $t = 0$

2. While **not converged**

- Compute the gradient

$$\nabla_{\theta} L_D(\theta^{(t)}) = 2X^T X \theta^{(t)} - 2X^T y$$

- Update the weights

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla_{\theta} L_D(\theta^{(t)})$$

- Increment  $t$ :

$$t = t + 1$$

- Output  $\theta^{(t)}$

# Gradient Descent for Linear Regression – Change Step Size

- Dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
- 1. Initialize  $\theta^{(0)} = 0$  (zero vector) and set  $t = 0$
- 2. While **not converged**
  - Compute the gradient

$$\nabla_{\theta} L_D(\theta^{(t)}) = 2X^T X \theta^{(t)} - 2X^T y$$

- Update the weights

$$\theta^{(t+1)} = \theta^{(t)} - \eta_t \nabla_{\theta} L_D(\theta^{(t)})$$

$$\eta_t = \frac{\eta_0}{n\sqrt{t+1}}$$

- Increment  $t$ :

$$t = t + 1$$

- Output  $\theta^{(t)}$

# Gradient Descent for Linear Regression – Change Step Size

- Dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$

1. Initialize  $\theta^{(0)} = 0$  (zero vector) and set  $t = 0$

2. While **not converged**

- Compute the gradient

$$\nabla_{\theta} L_D(\theta^{(t)}) = 2 \sum_{i=1}^n (\theta^{(t)T} x^{(i)} - y^{(i)}) x^{(i)}$$

- Update the weights

$$\theta^{(t+1)} = \theta^{(t)} - \eta_t \sum_{i=1}^n (\theta^{(t)T} x^{(i)} - y^{(i)}) x^{(i)} \quad \eta_t = \frac{\eta_0}{n\sqrt{t+1}}$$

- Increment  $t$ :

$$t = t + 1$$

- Output  $\theta^{(t)}$

# Distributed Computation of $\nabla_{\theta} L_D(\theta^{(t)})$

<b>Worker</b>	$\begin{bmatrix} \leftarrow & x^{(1)T} & \rightarrow \\ \leftarrow & x^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(2)T} & \rightarrow \\ \leftarrow & x^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(5)T} & \rightarrow \\ \leftarrow & x^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nd)$ distributed storage (total)	
<b>Map</b>	$\left(\theta^{(t)T} x^{(i)} - y^{(i)}\right) x^{(i)}$	$\left(\theta^{(t)T} x^{(i)} - y^{(i)}\right) x^{(i)}$	$\left(\theta^{(t)T} x^{(i)} - y^{(i)}\right) x^{(i)}$	$O(nd)$ distributed work (total)	$O(d)$ local storage
<b>Reduce</b>	$\theta^{(t+1)} = \theta^{(t)} - \frac{\eta_0}{n\sqrt{t+1}} \sum_{i=1}^n \left(\theta^{(t)T} x^{(i)} - y^{(i)}\right) x^{(i)}$			$O(d)$ local work	$O(d)$ local storage

Reducer send the latest weight vector to worker

# Gradient Descent


- Pros

- Easily parallelizable
- Each individual iteration is cheap
  - Can be further improved using stochastic or mini-batch GD
- Guaranteed to converge on convex objective functions

- Cons

- Potentially slow convergence
- Introduction of a hyperparameter
- Network communication in each iteration

# Communication Hierarchy



Perform parallel  
and in-memory  
computation  
whenever  
possible



# Minimize network communication

- Need to tradeoff between parallelism and network communication
- Three types of objects that may need to be communicated
  - Data
  - Model
  - Intermediate objects
- Strategies
  - Keep large objects local
  - Reduce the number of iterations

# Data Parallel: Compute outer products locally

<b>Worker</b>	$\begin{bmatrix} \leftarrow & x^{(1)T} & \rightarrow \\ \leftarrow & x^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(2)T} & \rightarrow \\ \leftarrow & x^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(5)T} & \rightarrow \\ \leftarrow & x^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nd)$ distributed storage (total)	
<b>Map</b>	$x^{(i)}x^{(i)T}$	$x^{(i)}x^{(i)T}$	$x^{(i)}x^{(i)T}$	$O(nd^2)$ distributed work (total)	$O(d^2)$ local storage
<b>Reduce</b>	$\left( \sum_{i=1}^n x^{(i)}x^{(i)T} \right)^{-1}$			$O(d^3)$ local work	$O(d^2)$ local storage

# Data Parallel: Compute pointwise gradients locally

<b>Worker</b>	$\begin{bmatrix} \leftarrow & x^{(1)T} & \rightarrow \\ \leftarrow & x^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(2)T} & \rightarrow \\ \leftarrow & x^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(5)T} & \rightarrow \\ \leftarrow & x^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nd)$ distributed storage (total)
<b>Map</b>	$\left(\theta^{(t)T} x^{(i)} - y^{(i)}\right) x^{(i)}$	$\left(\theta^{(t)T} x^{(i)} - y^{(i)}\right) x^{(i)}$	$\left(\theta^{(t)T} x^{(i)} - y^{(i)}\right) x^{(i)}$	$O(nd)$ $O(d)$ distributed      local work (total)      storage
<b>Reduce</b>	$\theta^{(t+1)} = \theta^{(t)} - \frac{\eta_0}{n\sqrt{t+1}} \sum_{i=1}^n \left(\theta^{(t)T} x^{(i)} - y^{(i)}\right) x^{(i)}$			$O(d)$ $O(d)$ local      local work      storage

Reducer send the latest weight vector to worker

# Model Parallel: Train each hyperparameter setting on different machine(s)

- Suppose we want to compare multiple hyperparameter settings  $\lambda_1, \lambda_2, \dots, \lambda_k$
- For  $i = 1, 2, \dots, k$ 
  - Train a model on  $D_{train}$  using  $\lambda_i$
- Evaluate each model on  $D_{val}$  and find the best hyperparameter setting,  $\lambda_{i^*}$
- Compute the error of a model trained with  $\lambda_{i^*}$  on  $D_{test}$



# Summary

1. Computation and storage should be linear in  $n$  and  $k$ 
  - For linear regression
    - When  $d$  is small, distribute matrix computation using outer products
    - When  $d$  is large, minimize squared error via distributed gradient descent
2. Perform parallel and in-memory computation whenever possible
3. Minimize network communication
  - Data vs model parallelism